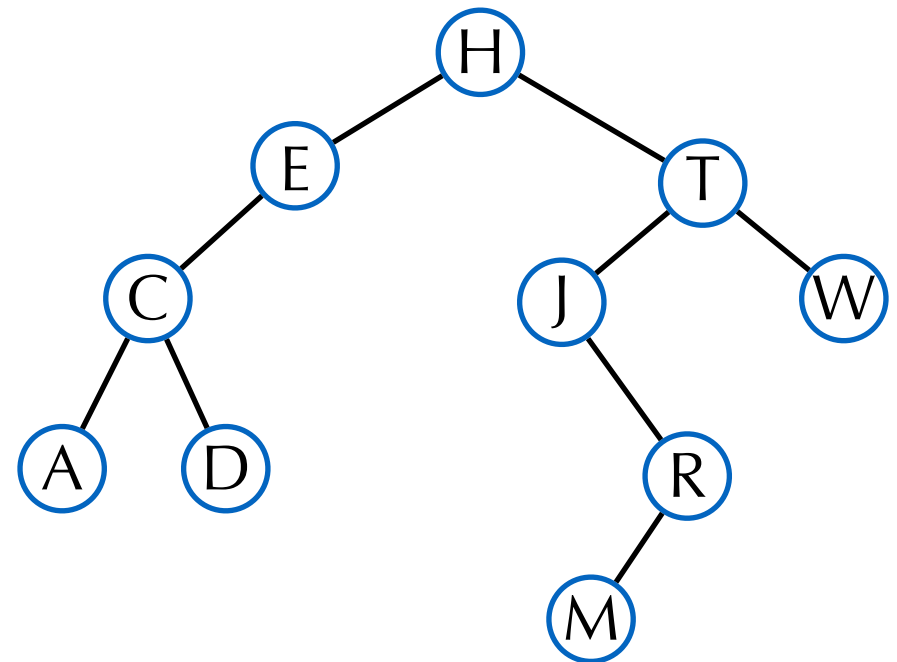


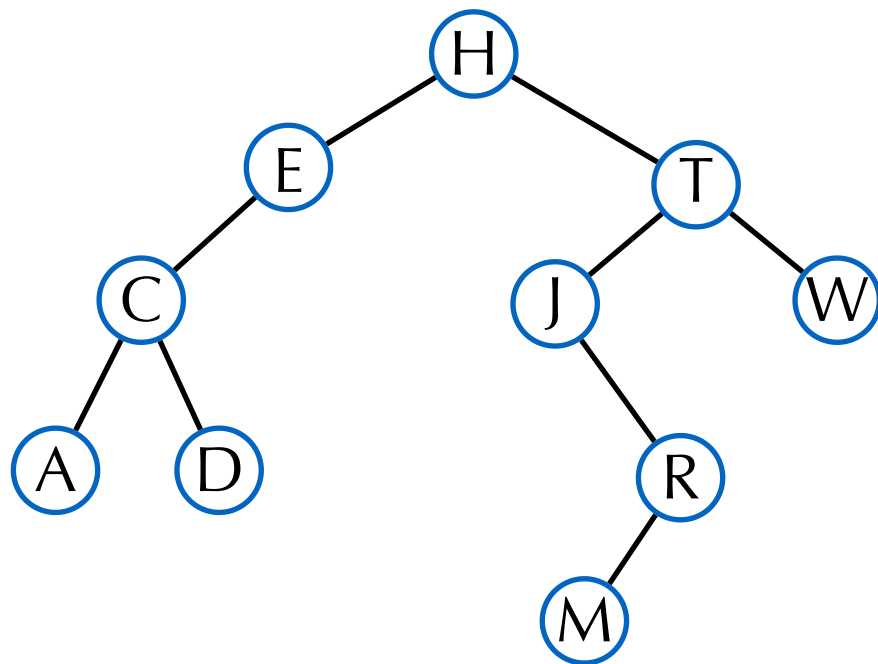
**CSE202**  
**Design and Analysis of Algorithms**

***Week 10 — Balance against Worst-Case***

## II. Binary Search Trees



# Recall Definition (CSE101 & 102)



Smaller elements to the left,  
larger elements to the right

```
class Node:
```

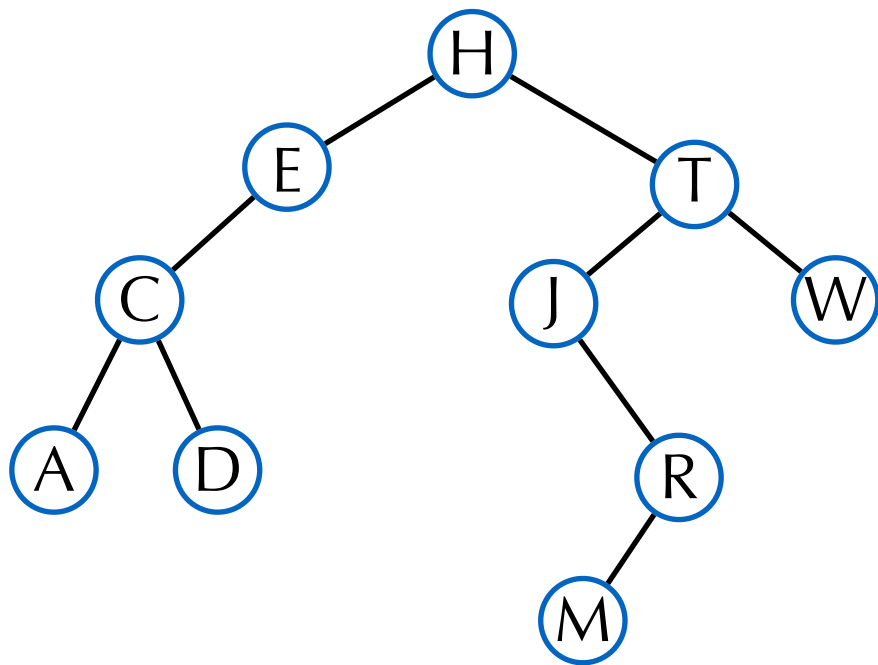
```
    def __init__(self, key, left=None, right=None):  
        self.key = key  
        self.left = left  
        self.right = right
```

```
class BST:
```

```
    def __init__(self):  
        self.root = None  
  
    def find(self, key):  
        return self._find(self.root, key)  
  
    def insert(self, key):  
        self.root = self._insert(self.root, key)  
  
    def delete(self, key):  
        self.root = self._delete(self.root, key)
```

# Find/Insert

```
def _find(self, node, key):  
    if node is None: return False  
    if node.key > key: return self._find(node.left, key)  
    if node.key < key: return self._find(node.right, key)  
    return True
```



```
def _insert(self, node, key):  
    if node is None: return Node(key)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    return node
```

Delete slightly more complicated (CSE102)

**Worst-case:** search in  $O(n)$  comparisons for a BST built from  $n$  keys.

# Average-Case Analysis

Internal path length:

$P_n :=$  sum depths of all nodes

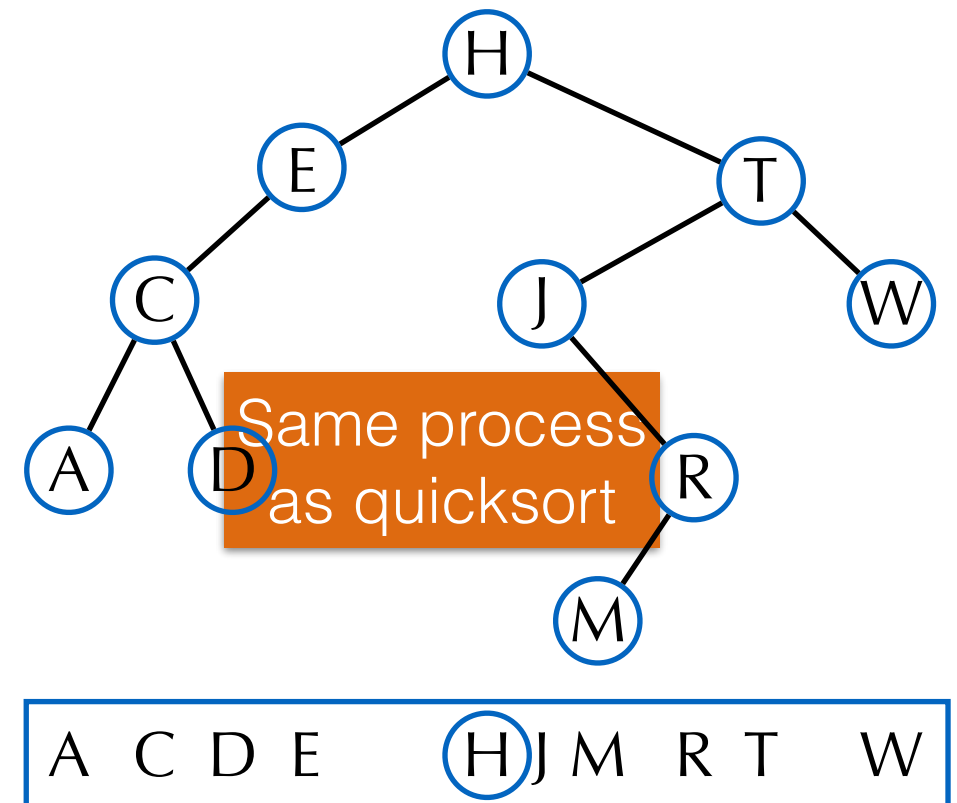
$P_n/n + 1$  : average successful search

$P_n/n + 3$  : average unsuccessful search  
(= insert)

Blackboard  
proof

**Prop.** In a BST built from  $n$  random keys, the average number of comparisons for a search is

$$1.39 \log_2 n + O(1)$$

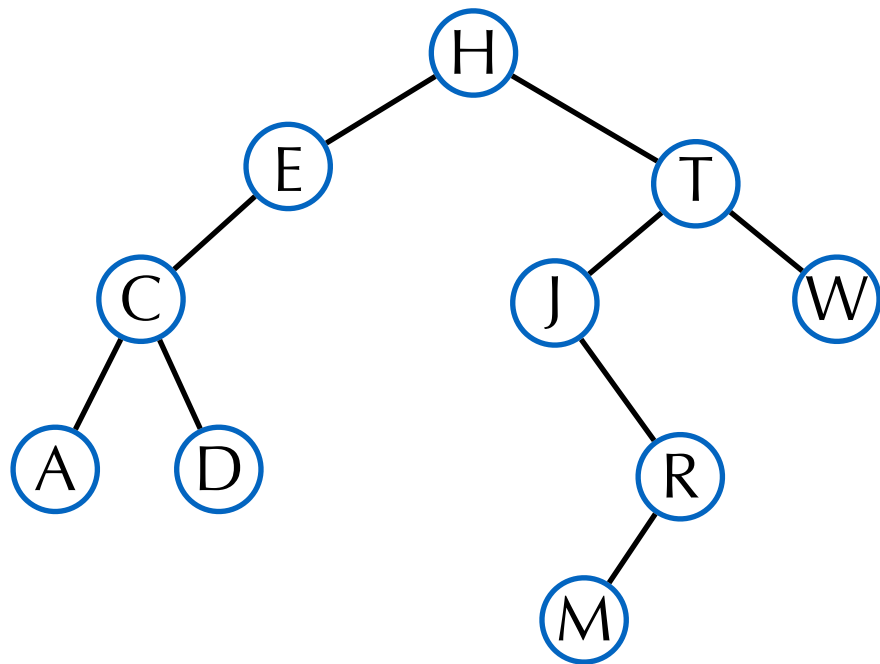


$$P_0 = P_1 = 0$$

$$\mathbb{E}P_n = n - 1 + \sum_{i=1}^n \frac{\mathbb{E}P_{i-1} + \mathbb{E}P_{n-i}}{n}$$

Same recurrence as in  
the analysis of quicksort.

# Select



min, max, floor, ceiling: easy

median, select:

change nodes into  
key, left, right, **size**

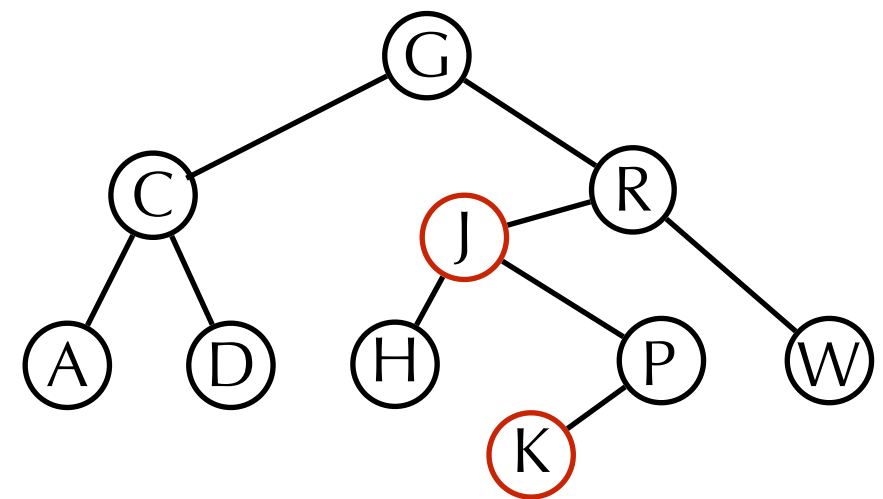
floor: largest key  
smaller than input

```
def _insert(self, node, key):  
    if node is None: return Node(key)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    node.size = 1 + size(node.left) + size(node.right)  
    return node
```

All these operations have cost bounded by the height,  
which is logarithmic **on average**.

Generalizes to  
higher dimensions  
(quadrees).

### III. Red-Black BST

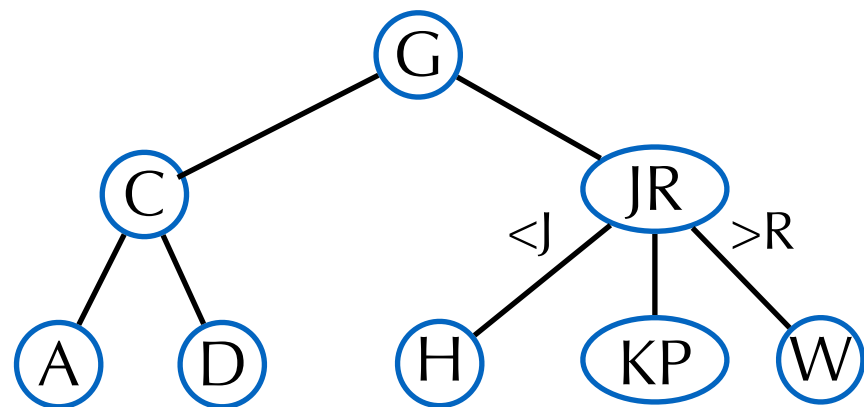


# WarmUp: 2-3 Search Trees

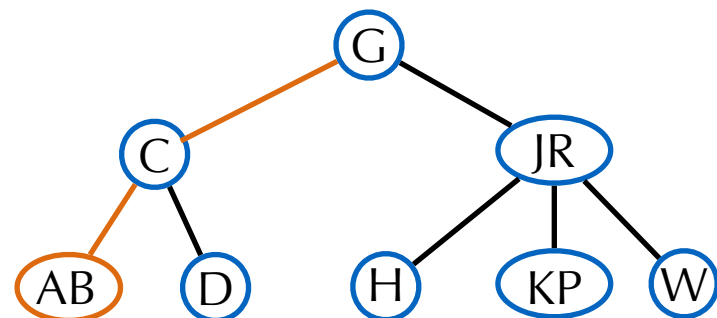
All leaves at the same level

Find: same as BST

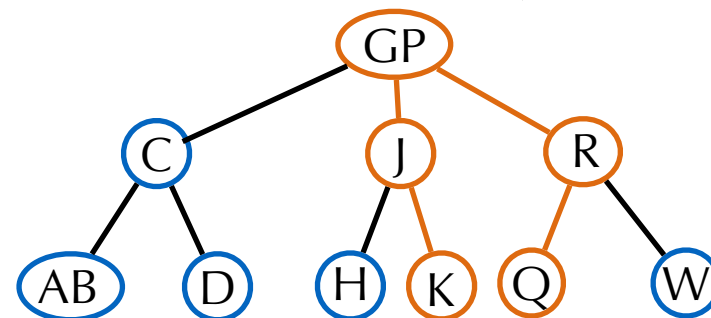
Insert maintaining perfect balance:  
search, insert at bottom  
and propagate upwards



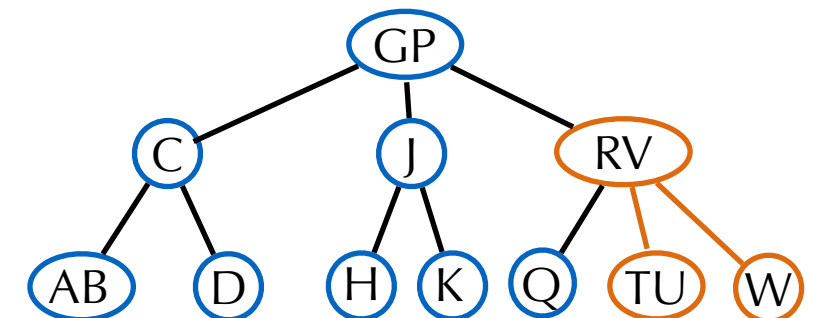
Insert B



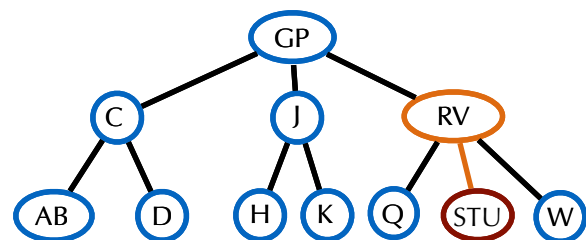
Insert Q



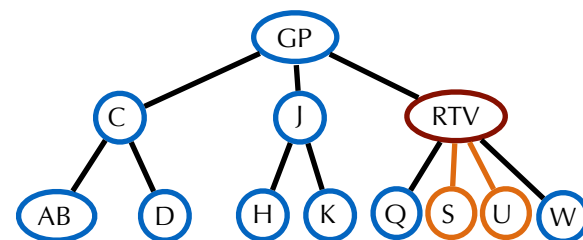
Insert V,U,T



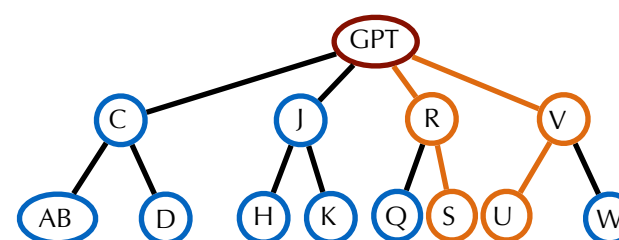
Insert S (1/4)



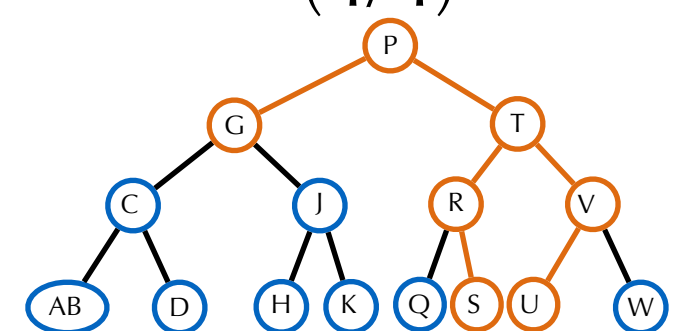
(2/4)



(3/4)



(4/4)

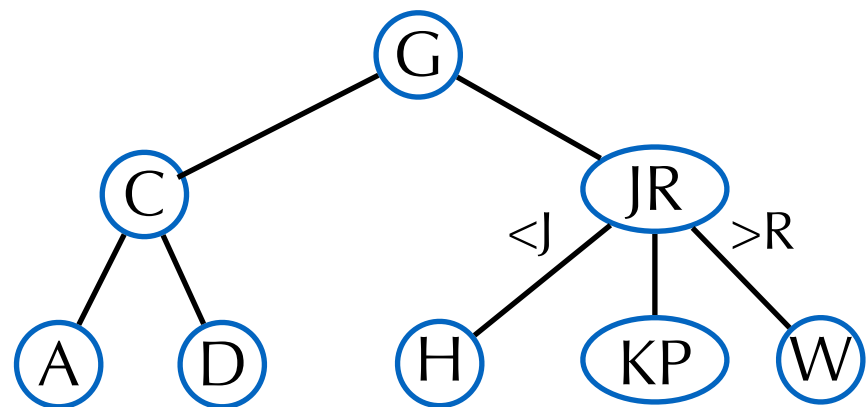


has to disappear

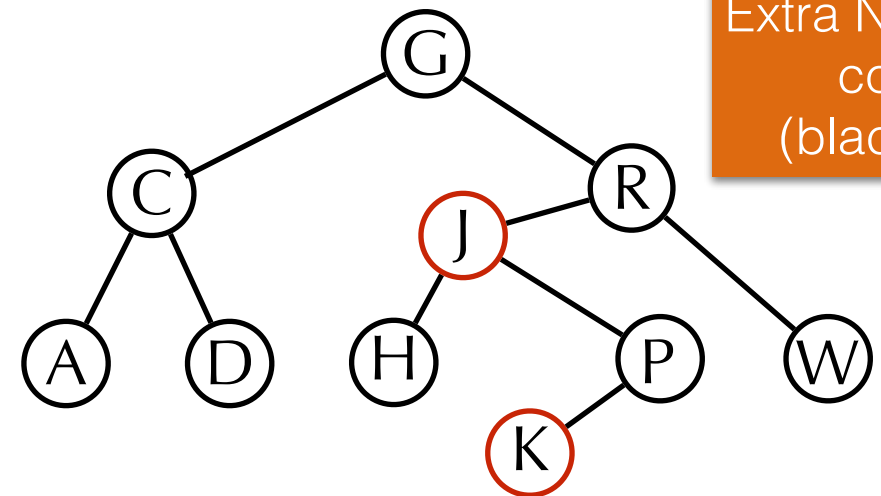
Worst-case number of nodes visited for find/insert:  $\log_2 n$



# (Left-Leaning) Red-Black Trees



stored as a  
coloured BST



Extra Node Info:  
colour  
(black root)

## Properties:

1. red nodes are left children;
2. red nodes have black children;
3. every path from the root to a leaf has the same number of black nodes.

Black  
balance

Red-black trees with these properties are in 1-to-1 correspondance with 2-3 trees.

find, select: code for BST unchanged! Just faster.

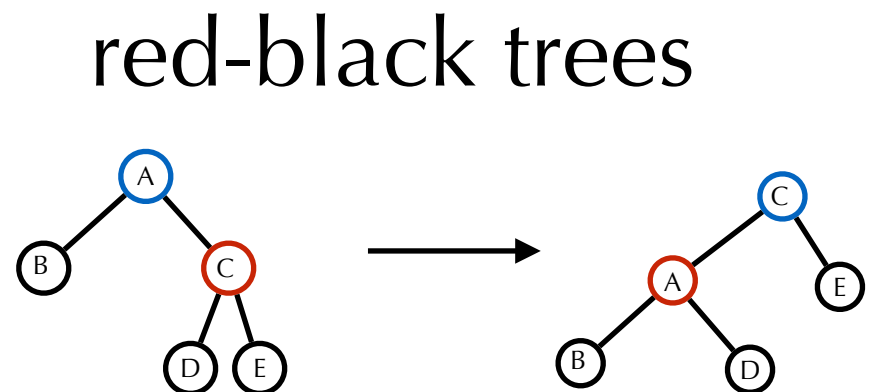
# Insertion

Insert maintaining  
order & black balance:  
search, insert **red** node at  
bottom and propagate upwards

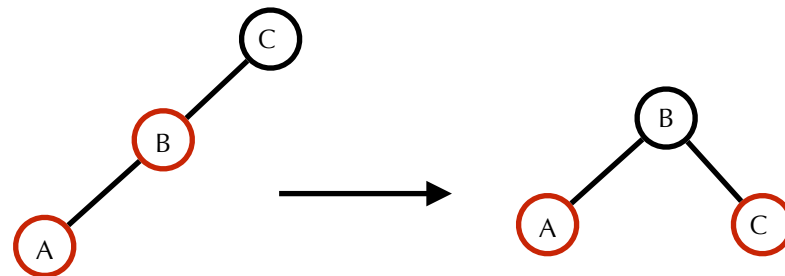
```
def _insert(self, node, key):  
    if node is None: return Node(key, red=True)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    if isRed(node.right) and not isRed(node.left): node = rotateleft(node)  
    if isRed(left.red) and isRed(node.left.left): node = rotateright(node)  
    if isRed(node.left) and isRed(node.right): flipcolors(node)  
    node.size = 1+size(node.left)+size(node.right)  
    return node
```

## Local fixes

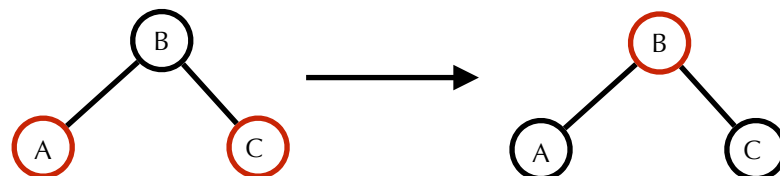
rotateleft



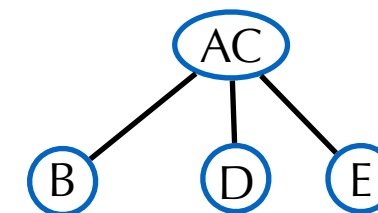
rotateright



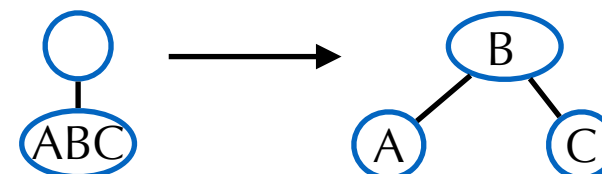
flipcolors



## 2-3 tree



ABC



Check that  
order &  
black balance  
are preserved

Delete more  
complicated

# Worst-Case Analysis

**Prop.** The height of a red-black BST with  $n$  nodes is bounded by  $2 \log_2 n$ .

Proof: exercise.

## Summary

algorithm (data structure)	worst-case cost (after N inserts)		average-case cost (after N random inserts)	
	search	insert	search hit	insert
<i>sequential search (unordered linked list)</i>	$N$	$N$	$N/2$	$N$
<i>binary search (ordered array)</i>	$\lg N$	$N$	$\lg N$	$N/2$
<i>binary tree search (BST)</i>	$N$	$N$	$1.39 \lg N$	$1.39 \lg N$
<i>2-3 tree search (red-black BST)</i>	$2 \lg N$	$2 \lg N$	$1.00 \lg N$	$1.00 \lg N$

Empirical.  
No proof yet.

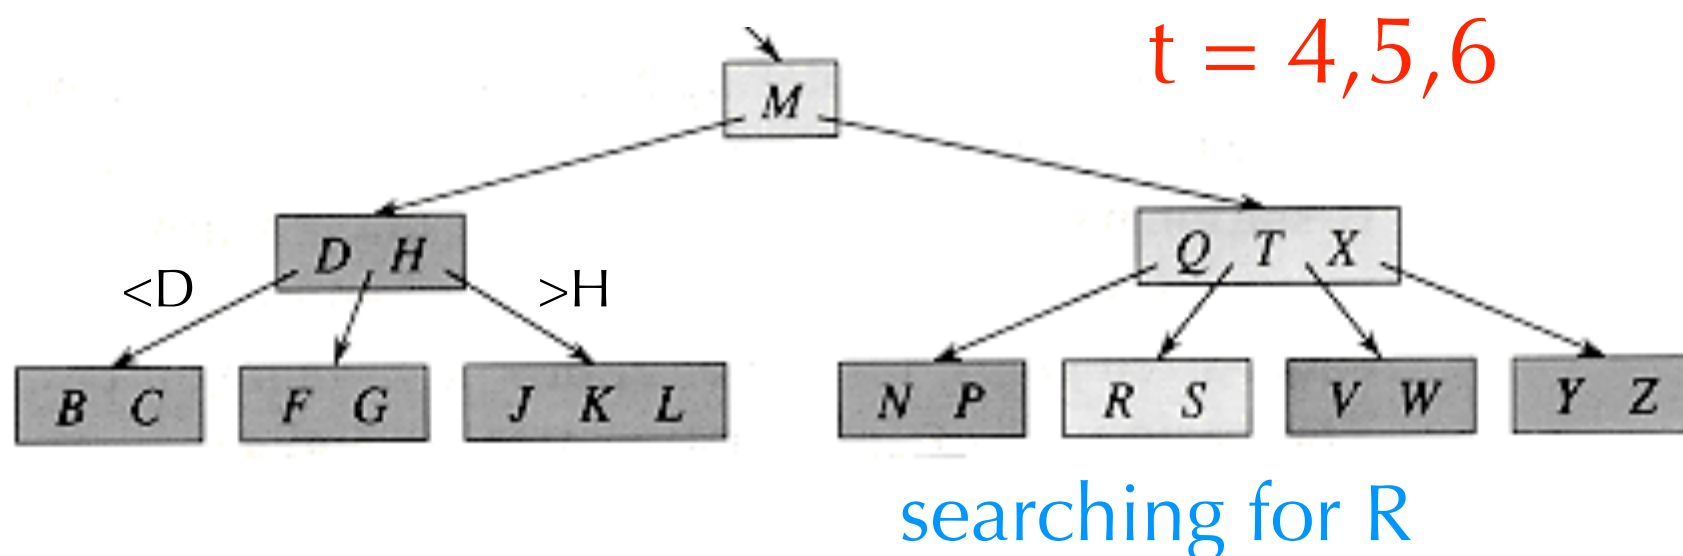
# I. B Trees

# B-Trees: Definition

Multi-way search trees commonly used in database systems (disk storage) - extension of 2-3 search trees

## Definition (B-tree of order $t \geq 3$ ):

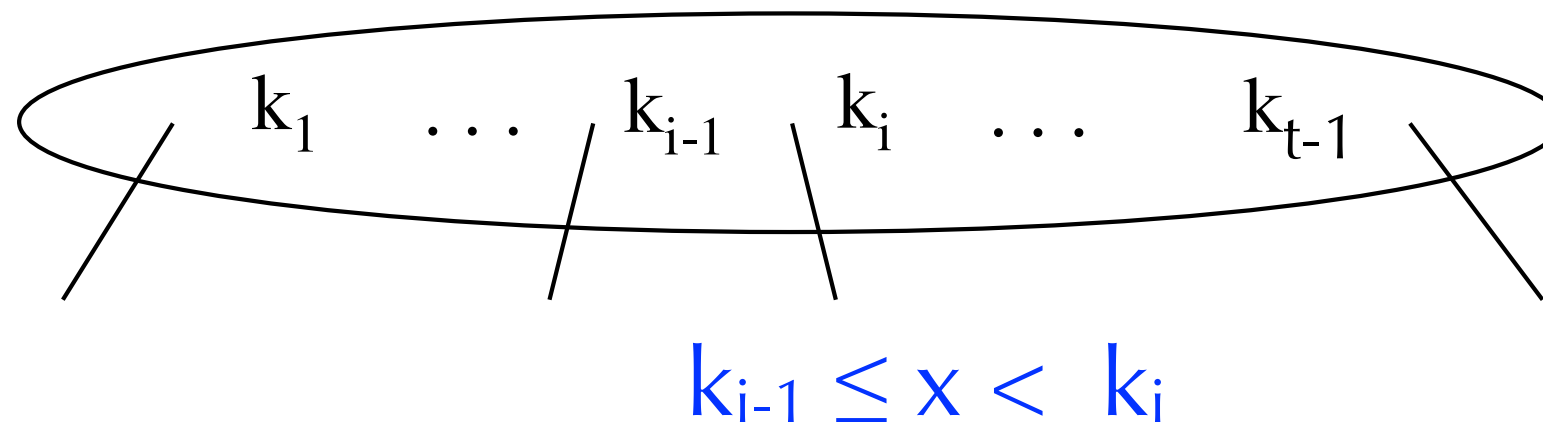
1. The root is either a leaf or has between 2 and  $t$  children;
2. Non-leaf nodes (except the root) have between  $\lceil t/2 \rceil$  and  $t$  children;
3. All leaves are at the same depth. Each leaf stores between  $\lceil t/2 \rceil$  and  $t$  items



# B-Trees: Internal Nodes

Each internal node of a B-tree has:

- between  $\lceil t/2 \rceil$  and  $t$  children
- up to  $t-1$  keys  $k_1 < k_2 < \dots < k_{t-1}$



Keys are ordered so that:

$$k_1 < k_2 < \dots < k_{t-1}$$

# B-Trees: Find

For a B-tree of order  $t$ :

- each internal node has up to  $t-1$  keys to search
- each internal node has between  $\lceil t/2 \rceil$  and  $t$  children
- depth of B-tree storing  $N$  items:  $O(\log_{\lceil t/2 \rceil} N)$

# B-Trees: Find

For a B-tree of order  $t$ :

- each internal node has up to  $t-1$  keys to search
- each internal node has between  $\lceil t/2 \rceil$  and  $t$  children
- depth of B-tree storing  $N$  items:  $O(\log_{\lceil t/2 \rceil} N)$

Complexity:

- $O(\log_2 t)$  to binary search which branch to take at a node
- Total time to find an item:

$$O(\log_2 t \cdot \log_{\lceil t/2 \rceil} N) = O(\log_2 N)$$



# B-Trees: Find

```
class Node:
```

```
    def __init__(self, leaf=False):  
        self.leaf = leaf  
        self.keys = []  
        self.children = []
```

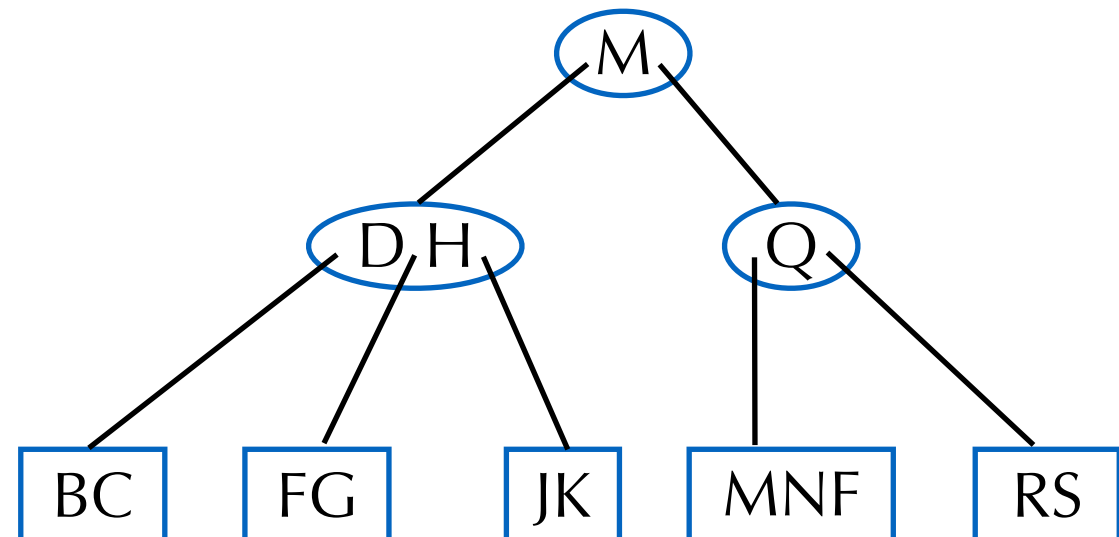
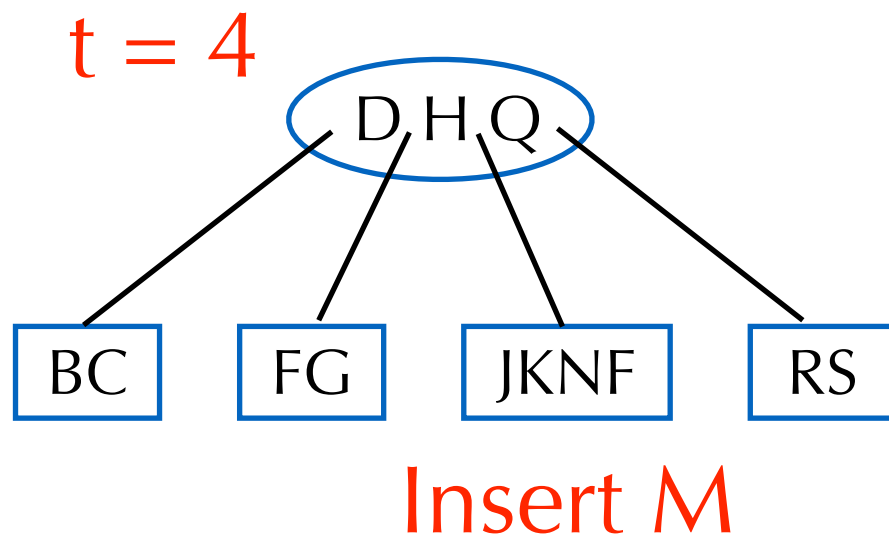
```
class BTree:
```

```
    def __init__(self, t):  
        self.root = Node(True)  
        self.t = t  
  
    def search(self, k):  
        x = self.root  
        i = 0  
        while i < len(x.keys) and k > x.keys[i][0]:  
            i += 1  
        if i < len(x.keys) and k == x.keys[i][0]:  
            return (x, i)  
        elif x.leaf:  
            return None  
        else:  
            return self.search(k, x.child[i])
```

# B-Trees: Insertion

**Insert x (similar to 2-3 search trees):** Do a find on x and find appropriate leaf node

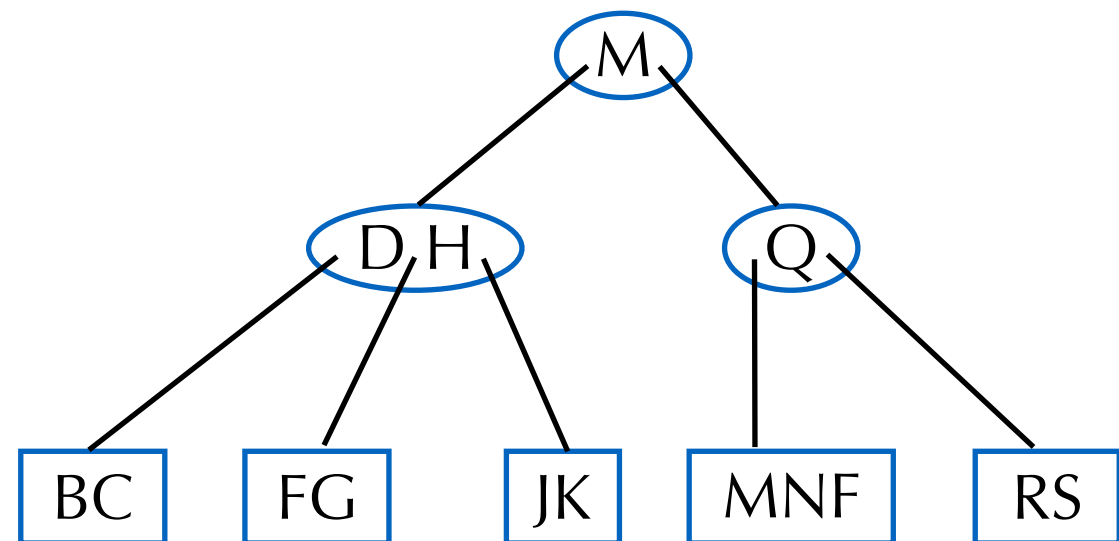
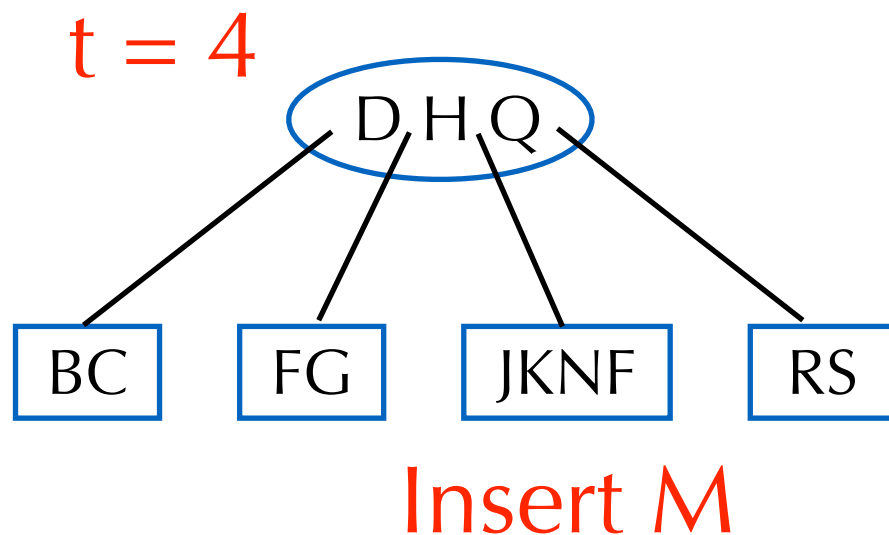
- if leaf node is not full, fill in empty slot with x
- if leaf node is full (has t items):
  - split into two nodes with  $\lfloor (t+1)/2 \rfloor$  and  $\lceil (t+1)/2 \rceil$  children
  - adjust parents up to the root node



# B-Trees: Insertion

**Insert x (similar to 2-3 search trees):** Do a find on x and find appropriate leaf node

- if leaf node is not full, fill in empty slot with x
- if leaf node is full (has t items):
  - split into two nodes with  $\lfloor (t+1)/2 \rfloor$  and  $\lceil (t+1)/2 \rceil$  children
  - adjust parents up to the root node



$$O(t \cdot \log_{\lceil t/2 \rceil} N) = O((t/\log_2 t) \cdot \log_2 N)$$

# B-Trees: Insertion

```
class Node:
```

```
    def __init__(self, leaf=False):  
        self.leaf = leaf  
        self.keys = []  
        self.children = []
```

```
def insert(self, k):  
    root = self.root  
    if len(root.keys) == self.t - 1:  
        temp = Node()  
        self.root = temp  
        temp.child.insert(0, root)  
        self.split_child(temp, 0)  
        self.insert_non_full(temp, k)  
    else:  
        self.insert_non_full(root, k)
```

```
def insert_non_full(self, x, k):  
    i = len(x.keys) - 1  
    if x.leaf:  
        x.keys.append((None, None))  
        while i >= 0 and k[0] < x.keys[i][0]:  
            x.keys[i + 1] = x.keys[i]  
            i -= 1  
        x.keys[i + 1] = k  
    else:  
        while i >= 0 and k[0] < x.keys[i][0]:  
            i -= 1  
        i += 1  
        if len(x.child[i].keys) == self.t - 1:  
            self.split_child(x, i)  
            if k[0] > x.keys[i][0]:  
                i += 1  
            self.insert_non_full(x.child[i], k)
```

# B-Trees: Insertion

```
class Node:
```

```
    def __init__(self, leaf=False):  
        self.leaf = leaf  
        self.keys = []  
        self.children = []
```

```
def insert(self, k):  
    root = self.root  
    if len(root.keys) == self.t - 1:  
        temp = Node()  
        self.root = temp  
        temp.child.insert(0, root)  
        self.split_child(temp, 0)  
        self.insert_non_full(temp, k)  
    else:  
        self.insert_non_full(root, k)
```

```
def split_child(self, x, i):  
    t = self.t  
    y = x.child[i]  
    z = Node(y.leaf)  
    x.child.insert(i + 1, z)  
    x.keys.insert(i, y.keys[ceil(t/2)])  
    z.keys = y.keys[ceil(t/2): t]  
    y.keys = y.keys[0: ceil(t/2)]  
    if not y.leaf:  
        z.child = y.child[ceil(t/2): t]  
        y.child = y.child[0: ceil(t/2)]
```

# B-Trees: Order Values

Tree in internal memory:  $t = 3$  or  $4$

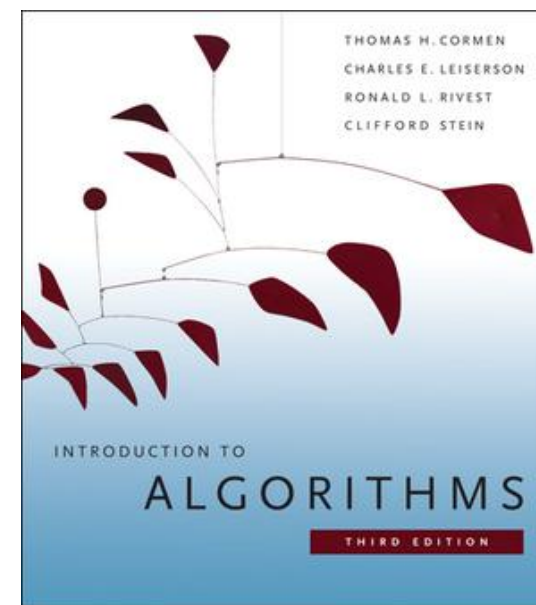
Tree on disk:  $t = 32$  to  $256$  (interior and leaf nodes fit on 1 disk block)

- depth = 2 or 3  $\rightarrow$  fast access to data in databases

# References for this lecture

The slides are designed to be self-contained.

They were prepared using the following book that I recommend if you want to learn more:



# Next

**NO** Assignment

Next tutorial: K-Dimensional Trees

Next week: String Algorithms 1 (Video)



# Feedback

Moodle

Questions: [constantin.enea@polytechnique.edu](mailto:constantin.enea@polytechnique.edu)